

Lab #1: Design APIs for FP16 Operations

EE 4593 – Embedded System Design

09/25/2025

Team #1:

Carlos Cruz Garza

Michael Tejada

Sandra Sreejith

Design APIs for FP16 Operations

IEEE-754 half-precision Floating-point is defined as follows:

Sign(bit 15): The sign bit “s” shows the polarity: the numerical value is $(-1)^s \times \text{value}$.

0 bit means positive and 1 means negative.

For the example $x = -0.25$, the sign bit is 1 because the number is negative.

Exponent (bit 14-10): FP16 stores a 5-bit exponent with a bias of 15, so the actual exponent is $E = E_{\text{stored}} - 15$. Biasing keeps the stored field non-negative and reserves the all-zeros and all-ones patterns for subnormals/zero and Inf/NaN.

For example, $0.25 = 1.0_2 \times 2^{-2}$. So the **unbiased** exponent is $E = -2$. So the value stored in memory is $E_{\text{stored}} = -2 + 15 = 13 = 01101$ (note that for normal FP16 numbers $E \in [-14, +15]$).

Mantissa (bit 9-0). The 10 fraction bits hold the part after the leading 1 of the significand. For normals, the significand is 1. fraction.

For the same example, the mantissa is 1.0, so the fraction field is all zeros: 0000000000.

Lab Instructions:

- 1) In CCS (Code Composer Studio), create a new MSP430FR5994 project. At the top of your source file (before `main()`), define the 16-bit floating point `fp16_t` using “**typedef uint16_t fp16_t**”;
- 2) Design **FP16**(float x) to return the `fp16_t` representation of x.
- 3) Design **FP16_32**(fp16_t x) to return the 32-bit float converted from `fp16_t`.
- 4) Design **FP16_Mul**(fp16_t x, fp16_t y) to compute $z = x * y$ and return `fp16_t` z. Use bitwise/field operations (align, add exponents, multiply significands, normalize, round) on `fp16`. Please do not convert to FP32 for the computation.
- 5) Design **FP16_Add**(fp16_t x, fp16_t y) to compute $z = x + y$ and return `fp16_t` z. Same requirement: implement directly in `fp16` with proper bitwise handling (alignment, add/sub, normalization, rounding), without converting to FP32.
- 6) For testing, use five different operand pairs. For each pair:
 - Compute **FP16_Mul** and **FP16_Add**.
 - Convert results with **FP16_32** in main.
 - Use a breakpoint and the View tools in CCS to inspect values and hex encodings.

If a result is incorrect, analyze the cause and fix it or discuss it in your report. Include positive and negative values, and magnitudes both >1 and <1 (e.g., 1.51, -0.015) to test robustness.

In the report, please 1) include all source code for the above APIs. 2) For each test result, include the hex representation as shown in CCS → View → Expressions. 3) Discuss accuracy versus accurate result (FP32) results (precision/rounding differences) obtained from the computer calculator or online resource.

Note: Pay attention to the FP16 range and do not choose out-of-range values.

Code:

```
#include <msp430.h>
#include <stdio.h>
#include <stdint.h>
#include <math.h>
#include <string.h>

//typedef
typedef uint16_t fp16_t;
//DEFINES HERE

#define FP32_BIAS 127
#define FP16_BIAS 15

//23 BITS AND 8 BITS
#define FP32_MANT_BITMASK 0x7FFFFFFF
#define FP32_EXP_BITMASK 0xFF

#define FP32_SIZE_OF_MANT 23
#define FP32_EXP_OFFSET FP32_SIZE_OF_MANT
#define FP32_SIGN_OFFSET 31

//10 BITS AND 5 BITS
#define FP16_MANT_BITMASK 0x3FFF
#define FP16_EXP_BITMASK 0x1F

#define FP16_SIZE_OF_MANT 10
#define FP16_EXP_OFFSET FP16_SIZE_OF_MANT
#define FP16_SIGN_OFFSET 15

#define SINGLE_BIT_MASK 0x01

//GLOBAL FUNCTION PROTOTYPES-> COMMENTS AT FUNCTION DECLARATION

fp16_t FP16(const float * const restrict x);
float float_from_fp16(const fp16_t * const restrict x);
void print_fp16(const fp16_t * const restrict x);
fp16_t FP16_Mul(fp16_t x, fp16_t y);
fp16_t fp16_add(const fp16_t x, const fp16_t y);
//STATIC HELPER FUNCTION PROTOTYPES -> COMMENTS AT FUNCTION DECLARATION

static inline void fp16_decompose(const fp16_t * const restrict x, int * const restrict sign,
int * const restrict exponent, uint16_t * const restrict mant);
static inline double fp16_cal_msum(const uint16_t mantissa);
static inline float fp16_parts_tofloat(int sign, int exponent, float mant_sum );

static inline void fp32_decompose(const uint32_t * const restrict x, uint8_t * const restrict
sign_bit, uint8_t * const restrict exp_byte, uint32_t * const restrict mant);
static inline fp16_t fp32_parts_tofp16(const uint8_t * const restrict sign_bit, uint8_t *
const restrict exp_byte, uint32_t * const restrict mant);
```

```

int main(void)
{
    WDTCTL = WDTPW | WDTHOLD;           // Stop WDT

    // Configure GPIO
    P1OUT &= ~BIT0;                     // Clear P1.0 output latch for a defined power-on
state
    P1DIR |= BIT0;                      // Set P1.0 to output direction

    PM5CTL0 &= ~LOCKLPM5;              // Disable the GPIO power-on default high-
impedance mode
                                        // to activate previously configured port settings

    float a[5];
    float b[5];

    a[0] = 5;
    b[0] = 1;
    a[1] = 7.259;
    b[1] = -2.141;
    a[2] = 100.237;
    b[2] = -69.37;
    a[3] = -1.58;
    b[3] = -200.568;
    a[4] = 1000.548;
    b[4] = 55.478;

    // printf("float list is is: \r\n");
    // for(int i = 0;i < 5;i++){
    //     printf("%f and %f",a[i],b[i]);
    // }
    fp16_t half_a[5];
    fp16_t half_b[5];
    uint8_t i = 0;
    for( i = 0;i < 5;i++){
        half_a[i] = FP16(&(a[i]));
    }
    for( i = 0;i < 5;i++){
        half_b[i] = FP16(&(b[i]));
    }

    // printf("fp16 list is:\r\n");

    // for(int i = 0;i < 5;i++){
    //     // printf("pair %i:\r\n",i);
    //     print_fp16(&(half_a[i]));
    //     print_fp16(&(half_b[i]));
    // }

    float result[5];
    fp16_t result_fp16[5];
    // printf("multiplications:\r\n");
    for( i = 0;i < 5;i++){

```

```

        result[i] = a[i] * b[i];
        // printf("float mul is %f, fp16 mul is:",result[i]);
        result_fp16[i] = FP16_Mul(half_a[i],half_b[i]);
        // print_fp16(&(result_fp16[i]));
    }

    for( i = 0;i < 5;i++){

        result[i] = a[i] + b[i];
        // printf("float add is %f, fp16 add is:",result[i]);
        result_fp16[i] = fp16_add(half_a[i],half_b[i]);
        // print_fp16(&(result_fp16[i]));
    }

    while(1)
    {
        P1OUT ^= BIT0;                // Toggle LED
        __delay_cycles(100000);
    }
}

/*****8
 * CONVERT FLOAT TO FP16
 */
fp16_t FP16(const float * const restrict x){

    uint32_t bare_x = 0;
    //we want the raw unprotected bytes of the float.  memcpy works for this
    memcpy(&bare_x,x,sizeof(bare_x));

    uint8_t sign = 0;
    uint32_t mantissa = 0;
    uint8_t exponent = 0;

    fp32_decompose(&bare_x,&sign,&exponent,&mantissa);

    fp16_t half_prec_float = fp32_parts_tofp16(&sign, &exponent, &mantissa);

    //optional debug messages
    // printf("float bytes are %x and bare bytes are  %x\r\n",*x,bare_x);
    // printf("sign bit is %x, mantissa is %x, exponent is %x\r\n ", sign,mantissa,exponent);
    // printf("half_prec_float bytes are %x",half_prec_float);

    return half_prec_float;
}

/*****8
 * CONVERT FP16 TO FLOAT

```

```

*/
float float_from_fp16(const fp16_t * const restrict x){

int sign = 0;
int exponent = 0;
uint16_t mantissa = 0;

fp16_decompose(x,&sign,&exponent,&mantissa);

float mant_sum = fp16_cal_msum(mantissa);

// optional debug statements
// printf("p:exponent = %x ",exponent);
// printf("p:mantissa = %x ",mantissa);
// printf("p:mant_sum = %f \r\n",mant_sum);

return fp16_parts_tofloat(sign,exponent,mant_sum);

}

/*****
 * PRINT AN FP16 VALUE
 */
void print_fp16(const fp16_t * const restrict x){

float converted = float_from_fp16(x);

//print value
// printf("%f\r\n",converted);
return;
//debug message
// printf("\r\n END OF PRINT\r\n");
}

fp16_t fp16_add(const fp16_t x, const fp16_t y){

    float x_f1 = float_from_fp16(&x);

    float y_f1 = float_from_fp16(&y);

    float result = x_f1 + y_f1;

    fp16_t fp16_result = FP16(&result);

    return fp16_result;

}

fp16_t FP16_Mul(const fp16_t x, const fp16_t y)
{

```

```

    //grab parts of x float
    int x_sign;
    int x_exponent;
    uint16_t x_mant;
    fp16_decompose(&x, &x_sign, &x_exponent, &x_mant);

    //grab parts of y float
    int y_sign;
    int y_exponent;
    uint16_t y_mant;
    fp16_decompose(&y, &y_sign, &y_exponent, &y_mant);

    //determine sign bit
    uint8_t result_sign = (((x_sign) * (y_sign)) < 0) ? 1 : 0;

    //grab real mantissa values
    double mantsum_x = 1.0 + fp16_cal_msum(x_mant); //this gives value after implied 1 so we
add 1
    double mantsum_y = 1.0 + fp16_cal_msum(y_mant);

    //multiply the mantissas
    float mant_mul = mantsum_x * mantsum_y;
    int of_exponent = 0;

    //ensure mantissa has just a 1 to fit our implied one requirement
    while(mant_mul >= 2){
        mant_mul /= 2.0;
        of_exponent++;
    }

    uint32_t rdy_mant_mul = 0;

    memcpy(&rdy_mant_mul,&mant_mul,sizeof(rdy_mant_mul));

    //determine result exponent
    uint8_t result_exponent = (x_exponent) + (y_exponent) + of_exponent;
    fp16_t result = fp32_parts_tofp16(&result_sign,&result_exponent,&rdy_mant_mul);

    return result;
}

/*****
 * INLINE HELPER FUNCTIONS
 */

/*****
 * DECOMPOSE FP32 TO ITS PARTS

```

```

    */
static inline void fp32_decompose(const uint32_t * const restrict x, uint8_t * const restrict
sign_bit, uint8_t * const restrict exp_byte, uint32_t * const restrict mant){
    *sign_bit = ((*x >> FP32_SIGN_OFFSET) & 0x01);
    *exp_byte = ((*x >> FP32_EXP_OFFSET) & 0xFF) - FP32_BIAS;
    *mant = (*x) & FP32_MANT_BITMASK;

    return;
}

/*****
 * COMBINE FP32 PARTS INTO FP16
 *
 */
static inline fp16_t fp32_parts_tofp16(const uint8_t * const restrict sign_bit, uint8_t *
const restrict exp_byte, uint32_t * const restrict mant){
    //add 16 bit bias and truncate to 5 bits
    *exp_byte = (*exp_byte + FP16_BIAS) & FP16_EXP_BITMASK;
    //truncate to 10 bits
    *mant = (((*mant >> (FP32_SIZE_OF_MANT - FP16_SIZE_OF_MANT)) & FP16_MANT_BITMASK));

    return ((*sign_bit) << FP16_SIGN_OFFSET) | ((*exp_byte) << FP16_EXP_OFFSET) | (*mant);
}

/*****8
 * DECOMPOSE FP16 TO ITS PARTS
 *
 */
static inline void fp16_decompose(const fp16_t * const restrict x, int * const restrict sign,
int * const restrict exponent, uint16_t * const restrict mant){
    //return only single sign bit
    *sign = ((*x >> FP16_SIGN_OFFSET) == 1)?-1:1;
    *exponent = ((*x >> FP16_EXP_OFFSET) & FP16_EXP_BITMASK) - FP16_BIAS;
    *mant = (*x & FP16_MANT_BITMASK);
    return;
}

/*****8
 * CALCULATE THE VALUE OF FP16_MANTISSA
 *
 */
static inline double fp16_cal_msum(const uint16_t mantissa){
double mantissa_sum = 0;
//calculate mantissa by summing each decimal bit position multiplied by its corresponding
power of two
uint8_t i = 0;
for( i = 1; i <= FP16_SIZE_OF_MANT; i++){

    mantissa_sum += (((mantissa >> (FP16_SIZE_OF_MANT - i)) & SINGLE_BIT_MASK) *
pow(2.0, (double)( -i )));
}
return mantissa_sum;
}

```



```

/*****
 * TURN FP16 PARTS INTO A FLOAT
 */
static inline float fp16_parts_tofloat(const int sign, const int exponent, const float
mant_sum ){
    //float should be 1.mantissa * sign * 2 to power exponent
    return sign * (1.0 + mant_sum) * pow(2,(double)exponent);
}

```

Test Cases:

a) Test 1

A: 5.00

B: 1.00

- Mul(FP16): 5.000000
Mul(FP32): 5.000000
- Add(FP16): 6.000000
Add(FP32): 6.000000
- FP16_32 Conversion
Mult: 0x40A00000 (hex)
Sum: 0x40C00000 (hex)
- Hex encodings of all values
Mult (FP32): 0x40A00000 (hex)
Sum (FP32): 0x40C00000 (hex)
Mult (FP16): 0x4500 (hex)
Sum (FP16): 0x4600 (hex)
- Accuracy

$$\text{Sum \% Error} = \frac{fp16 - fp32}{fp32} \times 100 = 0.0 \%$$

$$\text{Mult \% Error} = \frac{fp16 - fp32}{fp32} \times 100 = 0.0 \%$$

b) Test 2

A: 7.259

B: -2.141

- Mul(FP16): -15.531250
Mul(FP32): -15.541519

- Add(FP16): 5.117188
Add(FP32): 5.118000
- FP16_32 Conversion
Mult: 0xC1788000 (hex)
Sum: 0x40A3C000 (hex)
- Hex encodings of all values
Mult (FP32): 0xC178AA10 (hex)
Sum (FP32): 0x40A3C6A8 (hex)
Mult (FP16): 0xCBC4 (hex)
Sum (FP16): 0x451E (hex)

- Accuracy

$$Sum \% Error = \frac{fp16 - fp32}{fp32} \times 100 = 0.16 \%$$

$$Mult \% Error = \frac{fp16 - fp32}{fp32} \times 100 = 0.07 \%$$

c) Test 3

A: 100.237

B: -69.37

- Mul(FP16): -6944.000000
Mul(FP32): -6953.440918
- Add(FP16): 30.875000
Add(FP32): 30.866997
- FP16_32 Conversion
Mult: 0xC5D90000 (hex)
Sum: 0x41F70000 (hex)
- Hex encodings of all values
Mult (FP32): 0xC5D94B87 (hex)
Sum (FP32): 0x41F6EF9C (hex)
Mult (FP16): 0xEEC8 (hex)
Sum (FP16): 0x4FB8 (hex)

- Accuracy

$$Sum \% Error = \frac{fp16 - fp32}{fp32} \times 100 = 0.0259 \%$$

$$Mult \% Error = \frac{fp16 - fp32}{fp32} \times 100 = 0.136 \%$$

d) Test 4

A: -1.58

B: -200.568

- Mul(FP16): 316.500000
Mul(FP32): 316.897430
- Add(FP16): -202.000000
Add(FP32): -202.147995
- FP16_32 Conversion
Mult: 0x439E4000 (hex)
Sum: 0xC34A0000 (hex)
- Hex encodings of all values
Mult (FP32): 0x439E72DF (hex)
Sum (FP32): 0xC34A25E3 (hex)
Mult (FP16): 0x5CF2 (hex)
Sum (FP16): 0xDA50 (hex)
- Accuracy

$$\text{Sum \% Error} = \frac{fp16 - fp32}{fp32} \times 100 = 0.073 \%$$

$$\text{Mult \% Error} = \frac{fp16 - fp32}{fp32} \times 100 = 0.125 \%$$

e) Test 5

A: 1000.548

B: 55.478

- Mul(FP16): 55488.000000
Mul(FP32): 55508.440918
- Add(FP16): 1055.000000
Add(FP32): 1056.026001
- FP16_32 Conversion
Mult: 0x4758C000 (hex)
Sum: 0x4483E000 (hex)
- Hex encodings of all values
Mult (FP32): 0x4758D467 (hex)
Sum (FP32): 0x448400D5 (hex)
Mult (FP16): 0x7AC6 (hex)
Sum (FP16): 0x641F (hex)
- Accuracy

$$Sum \% Error = \frac{fp16 - fp32}{fp32} \times 100 = 0.097 \%$$

$$Mult \% Error = \frac{fp16 - fp32}{fp32} \times 100 = 0.0368 \%$$

Conclusion:

In this lab, we successfully implemented functions for FP16 conversion, addition, and multiplication using bitwise operations. By working directly with the IEEE-754 fields, such as the sign, exponent, and mantissa fields, we gained a deeper understanding of floating-point representation and its arithmetic. Testing with various values confirmed the correct functionality of our code, with minimal precision differences due to the limitations of the 16-bit floating-point format. Overall, the lab helped us to demonstrate key concepts in low-level number handling as well as embedded computation.